

Berk Durmus

Istanbul, Turkey · hello@berkdurmus.com · berkdurmus.com · github.com/berkdurmus

I take ideas from zero to production. Right now I'm building AI features at a health tech company—LLM agents, Playwright-based monitoring, RAG pipelines. Before that, co-founded a travel startup and grew it to 7K paying users. I spend my weekends making browser extensions and open source dev tools. MS in Machine Learning from TU Wien.

PROJECTS

Openmation — Browser Recording & Replay

openmation.dev · github.com/openmation/openmation

- Fully open source Chrome extension that records browser sessions, replays them with AI-assisted element matching, and generates shareable links
- Also supports scheduled runs for synthetic monitoring. TypeScript, Chrome APIs, LLM integration

TraceStation — Playwright Test Debugger

github.com/berkdurmus/trace-station

- Feeds Playwright traces into an LLM to explain why tests failed and where to look in the codebase
- Born from frustration with cryptic test failures at my day job

OTelMatic — AI-Powered Observability

github.com/berkdurmus/otelmatic

- Hooks into OpenTelemetry traces to do automated root cause analysis and anomaly detection. Basically an AI that watches your distributed system and tells you what broke

PromptTuner — Self-Optimizing Prompt Agent

github.com/berkdurmus/prompttuner

- Agent that writes prompts, tests them against eval sets, and iterates until metrics improve. Saves hours of manual prompt tweaking

Scapture — Screenshot-to-Notes Extension

github.com/berkdurmus/scapture

- Chrome extension: take a screenshot, AI turns it into organized notes

EXPERIENCE

Center Health — Remote, US

Senior Product Engineer, AI · Oct 2022 – Present

- A major clinic deal was about to fall through because we didn't support Freestyle Libre CGMs. I had 24 hours. Built the API connector, shadow-compared against our Dexcom data for validation, wrapped it in feature flags, added monitoring and rollback scripts. Live by morning with zero regressions. That template now cuts new device integrations from weeks to days.
- Set up Playwright tests across our patient and clinic apps, running on cron for synthetic monitoring. Bugs reaching production dropped by half.
- Own the AI stack: LLM agents for clinic workflows, RAG for medical context, code generation tools. Added human-in-the-loop checkpoints for anything involving clinical decisions.
- Built a memory system so our LLM agents remember patient context across 10K+ conversations (90% retrieval accuracy). Custom embeddings + retrieval pipeline.
- Company went from 2K to 18K monthly paying customers. Personally reworked the signup funnel—MRR up 20% in 3 months.

Pulfy.com — Turkey

Co-founder & CTO · Jul 2021 – Oct 2022

- Two technical co-founders, that's it. We wore every hat—backend, frontend, infra, product, support. Built the whole thing from scratch: data pipeline scraping 10K+ destinations for events, prices, weather, visa requirements, flight costs.
 - Grew to 7,000 paying users. Learned a lot about what happens when your scraper gets blocked at 2am.
 - Built the recommendation engine from scratch—collaborative filtering combined with content-based signals. Users could find personalized travel suggestions in under 200ms.

Turing — Remote, US and EU

Software Engineer · Jan 2019 – Oct 2021

- Contract work for EU clients. Built data crawlers, ML-based image classification tools, recommendation systems, streaming pipelines, and i18n infrastructure.

Turing — Remote, US and EU

Machine Learning Engineer · Aug 2017 – Jan 2019

- Built NLP pipelines for document classification and entity extraction. One client had 50K+ legal documents that needed auto-categorization—got it to 94% accuracy.
 - Trained custom CNN models for e-commerce product image recognition. Handled edge cases like partial occlusions and varied lighting.
 - Built RNN-based sequence models for time-series forecasting and sentiment analysis. Deployed to production serving real-time predictions.

SKILLS

TypeScript, React, Redux, MobX, Node.js, GraphQL, PostgreSQL, Redis, MongoDB. Chrome extension development. LLM agents, RAG, embeddings, prompt engineering, LangChain. Playwright for testing and monitoring. Cloudflare Workers, AWS, Docker.

EDUCATION

Vienna University of Technology — MS Informatics, Machine Learning focus	2017–2019
Bilkent University — BS Electrical & Electronics Engineering	2013–2017